

■■ DeepFakeShield AI: Multimodal Forensic Platform

Technical Dissertation & Presentation Guide

Author	Vanshika Tangri (Student No: 2315843)
Course	AI-4-Creativity
Objective	Comprehensive detection of manipulated media using multimodal AI analysis.

1. Project Overview & Motivation

DeepFakeShield AI is an advanced forensic ecosystem designed to address the escalating threat of synthetic media. Unlike heuristic-based detectors, it leverages a **multimodal consensus mechanism** that analyzes the three primary dimensions of digital video:

- Visual Consistency (Spatial & Temporal)
- Acoustic Authenticity (Voice Synthesis & Cloning)
- Audio-Visual Alignment (Lip-Sync Synchronization)

The "Truth Score" Philosophy

The platform doesn't just output a binary "Fake" or "Real" label. It provides a **Forensic Confidence Score** based on calibrated probabilities across all modalities, allowing human investigators to verify the findings through an interactive evidence timeline.

2. Comprehensive Technology Stack

A. AI & Machine Learning Infrastructure

Component	Technology	Detail
Video Model	ViT-B/16	12 Transformer layers, 12 heads, 768 hidden dim.
Audio Model	Wav2Vec2	Self-supervised learning on raw waveforms for spoof detection.
Sync Model	SyncNet	CNN-based cross-correlation for lip-audio alignment.
Inference	PyTorch 2.0+	Accelerated with CUDA/ROCM where available.
Computer Vision	OpenCV & PIL	Optimized frame extraction and spatial preprocessing.

B. Scalable Backend Architecture

Component	Technology	Role
API Layer	FastAPI	High-concurrency async endpoints (Pydantic V2).
Task Queue	Celery + Redis	Distributed processing to handle heavy 4K video analysis.
Database	PostgreSQL	Relational storage for forensic metadata and user logs.
Authentication	JWT (OAuth2)	Secure stateless authentication for forensic analysts.

3. Advanced System Architecture

Distributed Processing Model

DeepFakeShield is designed to run in containerized environments, allowing the ML Worker to scale independently of the API server.

```
graph TD
    subgraph Client_Space [Client Layer]
        User([Forensic Analyst]) --> Frontend[Vanilla JS Glassmorphism UI]
    end

    subgraph Service_Mesh [Service Layer]
        Frontend <--> API[FastAPI Gateway]
        API <--> DB[(PostgreSQL)]
        API <--> Redis[(Redis Message Broker)]
    end

    subgraph Analysis_Engine [ML Processing Cluster]
        Redis <--> Worker[Celery Worker Nodes]
        Worker --> V_Eng[Video Engine - ViT]
        Worker --> A_Eng[Audio Engine - Wav2Vec2]
        Worker --> L_Eng[Sync Engine - SyncNet]
        V_Eng & A_Eng & L_Eng --> Fusion[Weighted Fusion Module]
    end

    subgraph Delivery [Output Layer]
        Fusion --> PDF[PDF Report Generator]
        PDF --> User
    end
```

4. The Science Behind Detection

Phase 1: Spatial & Temporal Feature Learning

Our **Vision Transformer (ViT)** model views a frame as a sequence of 16x16 patches. This allows the model to learn:

- **Boundary Inconsistencies:** Artifacts at the seam where a fake face meets a real body.
- **Subtle Texture Blurring:** Deepfake generators often struggle with skin pores and micro-expressions.
- **GAN Signatures:** Spectral patterns left by Generative Adversarial Networks.

Phase 2: Audio Spoof Detection

Using **Wav2Vec2**, the system analyzes the raw waveform to detect:

- **Formant Anomalies:** Variations in vocal tract resonance that differ from human speech.
- **Phase Inconsistencies:** Discontinuities in synthetic audio segments.
- **Silence Analysis:** Artifacts in the quiet segments of a voice-cloned clip.

Phase 3: Cross-Modal Alignment (Lip-Sync)

The system calculates the **Correlation Coefficient** between the visual movement of the lips and the acoustic phonemes. A mismatch in timing (Lip-Sync lag) is a high-confidence indicator of a "Deepfake Reenactment" or "Voice Over Dubbing."

5. Multimodal Fusion Logic

The final verdict is calculated using a **Confidence-Weighted Ensemble**:

Score = (Σ (ModalityScore_i × Confidence_i)) / (Σ Confidence_i)

Modality	Default Weight	Key Artifacts Detected
Video	45%	Eye blinking anomalies, facial jitter, boundary noise.
Audio	30%	Spectral gaps, robotic resonance, pitch-shifting artifacts.
Lip-Sync	25%	Audio-Visual lag, phoneme-viseme mismatch.

6. Process Lifecycle

```
sequenceDiagram
    participant U as User
    participant B as Backend
    participant R as Redis
    participant W as ML Worker

    U->>B: Upload Media (SHA-256 integrity check)
    B->>R: Push Analysis Task
    R->>W: Fetch Task
    Note over W: Multithreaded Processing
    W->>W: Frame Extraction & Normalization
    W->>W: GPU-Accelerated ViT Inference
    W->>W: Audio Spectral Analysis
    W->>W: Cross-Correlation (Lip-Sync)
    W->>B: Persist Results to PostgreSQL
    B->>U: Notify Progress (WebSockets/Polling)
    U->>B: Generate PDF Forensic Report
```

7. Reliability & Forensic Integrity

Anti-Adversarial Measures

- **Input Sanitization:** Verifies MIME types and file signatures to prevent polyglot file attacks.
- **Data Integrity:** Every analysis result is linked to the **SHA-256 hash** of the original file to ensure chain of custody.

- **Explainability:** The platform highlights the specific **Evidence Segments** where the manipulation was most likely detected.

Future Improvement Roadmap

Active Learning Loop: Retraining models automatically on newly discovered deepfake variations (e.g., Diffusion-based faces).

Blockchain Notarization: Registering authenticity certificates on a public ledger for immutable proof.

Real-time Stream Analysis: Extending the platform to monitor live video broadcasts for synthetic content injection.

■ 8. Conclusion

DeepFakeShield AI represents a significant step forward in the fight against digital misinformation. By combining state-of-the-art transformer architectures with traditional signal processing, it provides a comprehensive defense mechanism for the modern era.

DeepFakeShield AI: Protecting Truth in the Age of Synthetic Media